

Homework 4

Helen Ngo


November 1, 2016

Problem 1 Use sklearn's implementation of k -Nearest Neighbors for regression purposes. Find the best value of k using 10-fold Cross-Validation (CV).

- (a) Generate 1000 data points.
- (b) Use 10-fold CV to report the three best values of k -neighbors that yield the best CV E_{out} ; vary the values of k in the following range: $k = 1, 3, 5, \dots, 2 \lfloor \frac{N+1}{2} \rfloor - 1$.
- (c) Report the best CV E_{out} .

Solution:

- (a) The data was generated using the given genDataSet function, modified so that it only returned the input x and target y values. The “Noisy Y” data was used as the target y values:



This can be seen in the python document attached.

(b)

```
def bestK(self, folds = 10):
    # Reshape dataset for k-neighbors regression
    x = (self.X).reshape(-1, 1)
    y = (self.Y).reshape(-1, 1)
    # Housekeeping
    Eout = []
    N = len(x)
    maxk = int(2*np.floor((N*(folds-1)/folds+1)/2)-1)
    # Output Eout list for k-neighbor range
    for n_neighbors in range(1, maxk, 2):
        kf = KFold(n_splits = folds)
        kscore = []
        for train, test in kf.split(x):
            x_train, x_test = x[train], x[test]
            y_train, y_test = y[train], y[test]
            reg = neighbors.KNeighborsRegressor(
                n_neighbors, weights = 'distance')
            reg.fit(x_train, y_train)
            kscore.append(reg.score(x_test,
                                    y_test))
        Eout.append(sum(kscore)/len(kscore))
    return Eout
```

Using the python code above, a list of E_{out} for the given range of k values was found. The index i of the max value for the list of E_{out} was determined, then used to calculate the corresponding k . Note that

k	index
1	0
3	1
5	2
\vdots	\vdots
$2 \lfloor \frac{N+1}{2} \rfloor - 1$	$\frac{k-1}{2}$

therefore the k value can be found from the index using the formula

$$k = 2(i) + 1.$$

The index of the three best values of k -neighbors that yield the best CV E_{out} were found: $k = 289, 296$ and 303 .¹

(c) The best CV $E_{out} = 0.245024051068$ was given by $k = 289$.

Problem 2 Generate the dataset in the same way as in Problem 1. Repeat the experiment 100 times storing the best three k number of neighbors in every single trial, and at the end of all the trials plot a histogram of all the saved values.

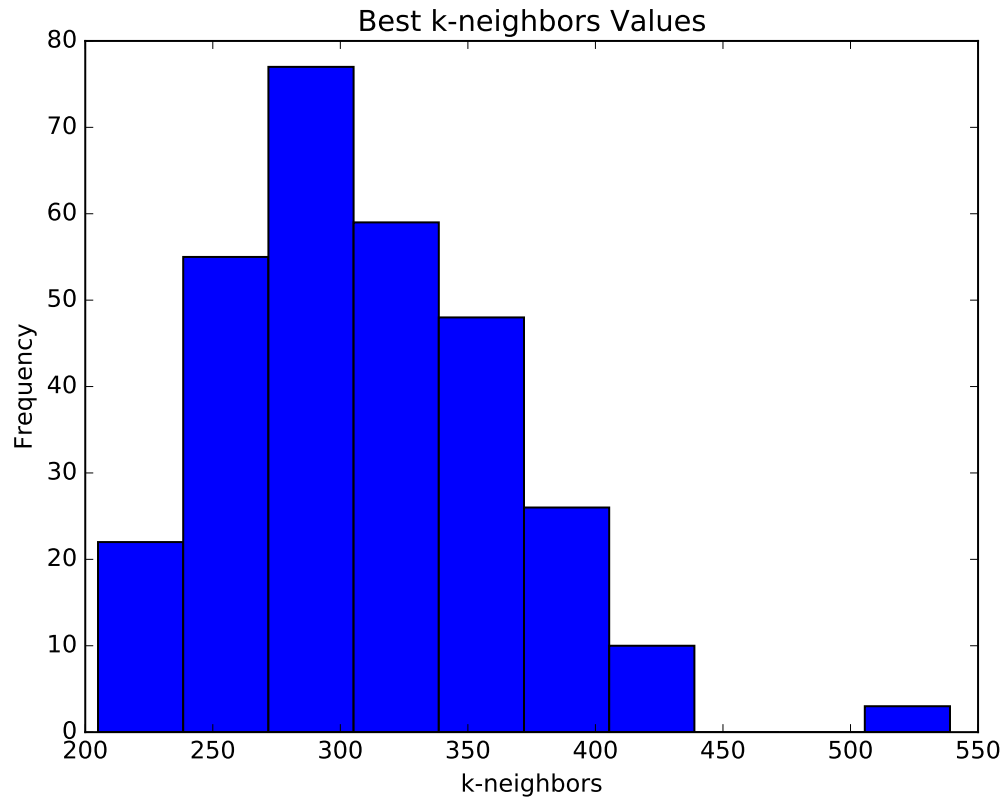
Solution: Commenting out the plotting feature of the `genDataSet` function, the experiment was repeated a 100 times.

```
def main():
    kvalues = np.empty([300,1]) # vector of best k's
    # Conducting 100 trials
    for trial in range(0,100):
        d = Neighbors(N=1000) # generate data
        Eout = d.bestK() # Eout array
        # Determining best k
        for j in range(0,3):
            i = Eout.index(max(Eout))
            k = 2*i+1
            kvalues[trial*3+j] = k
            print('k_{}_ '.format(j), k)
            print('Eout_{}_ '.format(j), max(Eout))
            del Eout[i]

    # Plotting histogram
    plt.hist(kvalues)
    plt.title("Best_k-neighbors_Values")
    plt.xlabel("k-neighbors")
    plt.ylabel("Frequency")
    plt.savefig('khistogram.pdf', bbox_inches='tight')
    plt.show()
```

¹Read footnote(2), these k values were properly adjusted.

The following histogram was generated using the best k values. ²



²The k number of neighbors may be 1 to 2 k values off due to method of finding them, but this will do for examining the trend, such as for a histogram.